

# ***Operating Systems***

**CSE 231**

**Instructor: Sambuddho Chakravarty**

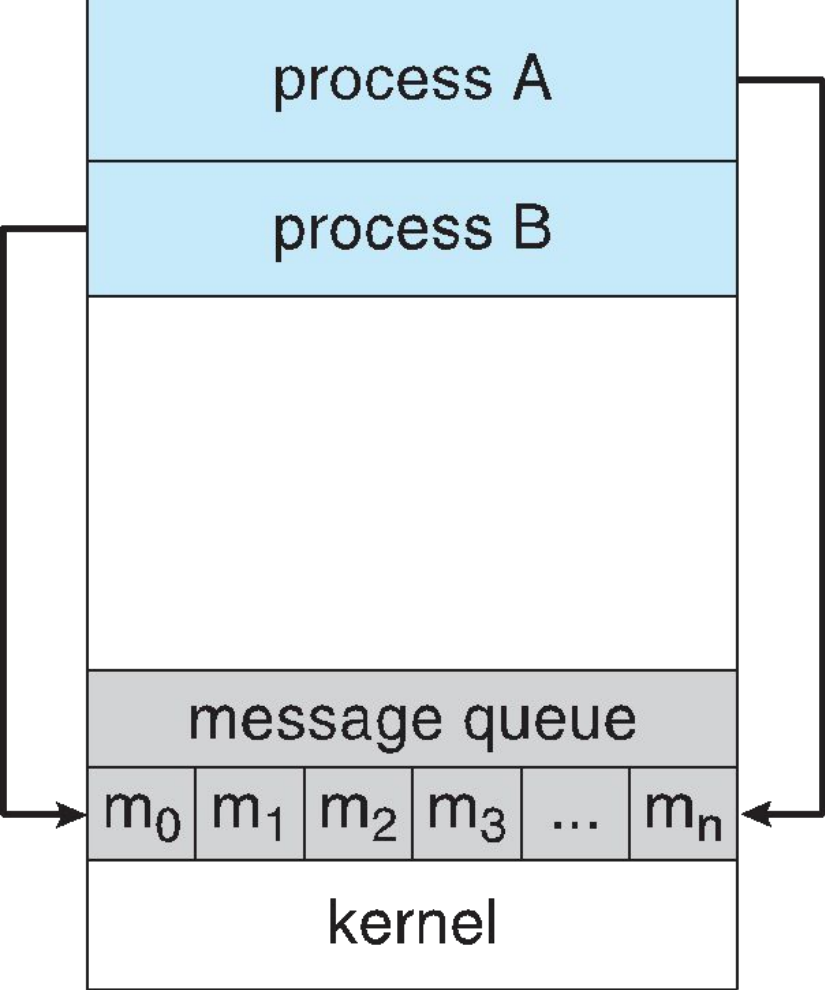
(Semester: Winter 2018)

Week 5: Jan 29 – Feb 2

# Interprocess Communication


- Processes within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
  - **Shared memory**
  - **Message passing**

# Communications Models



(a)

Message Passing



(b)

Shared Memory

# Message Passing v.s. Shared Memory

- Message passing
  - Why good? Simpler. All sharing is explicit
  - Why bad? Overhead. Data copying, cross protection domains
- Shared Memory
  - Why good? Performance. Set up shared memory once, then access w/o crossing protection domains
  - Why bad? Synchronization

# Signals

What is signal?

- A signal is an event generated by the UNIX system in response to some condition, upon which a process may in turn take some action.
- Signals are generated by some error conditions, such as memory segment violations, floating point processor errors or il-legal instructions. They are generated by the shell and terminal handlers to cause interrupts.

# Signals

Generally, signals can be generated, caught and acted upon, or ignored.

- Signal names are defined by including the header file `< signal.h >` as followed:
- SIGABORT: Process abort;
- SIGALRM: Alarm clock;
- SIGFPE: Floating point exception
- SIGHUP: Hangup
- SIGILL: illegal instruction
- SIGINT: Terminal interrupt.
- SIGKILL: kill (can't be caught or ignored)
- SIGPIPE: write on a pipe with no reader
- SIGQUIT: Termination quit
- SIGSEGV: Invalid memory segment access
- SIGTERM: Termination
- SIGUSR1: User-defined signal 1
- SIGUSR2: User-defined signal 2

# Signal Concepts

- Signals are defined in `<signal.h>`
- `man 7 signal` for complete list of signals and their numeric values
- `kill -l` for full list of signals on a system
  - 64 signals. The first 32 are traditional signals, the rest are for real time applications

# Signals

- Programming interface of signal handling

```
#include <signal.h>
void (*signal(int sig, void(*func)(int))) ;
```

It says that signal is a function that takes two parameters – sig and func



# POSIX Signal Handling

## C90 standard

- Defines **signal()** and **raise()** functions
  - Work across all systems (UNIX, LINUX, Windows), but...
  - Work **differently** across some systems!!!
    - On some systems, signals are blocked during execution of handler for that type of signal -- but not so on other (older) systems
    - On some (older) systems, handler installation for signals of type x is cancelled after first signal of type x is received; must reinstall the handler -- but not so on other systems
- Does not provide mechanism to block signals in general

# POSIX Signal Handling

## POSIX standard

- Defines **kill()**, **sigprocmask()**, and **sigaction()** functions
  - Work the same across all POSIX-compliant UNIX systems (Linux, Solaris, etc.), but...
  - Do not work on non-UNIX systems (e.g. Windows)
- Provides mechanism to block signals in general

# Blocking Signals in General

Each process has a signal mask in the kernel

- OS uses the mask to decide which signals to deliver
- User program can modify mask with `sigprocmask()`

`sigprocmask()`

```
int sigprocmask(int iHow,  
                const sigset_t *psSet,  
                sigset_t *psOldSet);
```

- `psSet`: Pointer to a signal set
- `psOldSet`: (Irrelevant for our purposes)
- `iHow`: How to modify the signal mask
  - `SIG_BLOCK`: Add `psSet` to the current mask
  - `SIG_UNBLOCK`: Remove `psSet` from the current mask
  - `SIG_SETMASK`: Install `psSet` as the signal mask
- Returns 0 iff successful

Functions for constructing signal sets

- `sigemptyset()`, `sigaddset()`, ...

# Blocking Signals Example

```
sigset_t sSet;
int main(void) {
    int iRet;
    sigemptyset(&sSet);
    sigaddset(&sSet, SIGINT);
    sigprocmask(SIG_BLOCK, &sSet, NULL);
    ...
    sigprocmask(SIG_UNBLOCK, &sSet, NULL);
    assert(iRet == 0);
    ...
}
```

# Installing a Signal Handler

## **sigaction()**

```
int sigaction(int iSig,  
              const struct sigaction *psAction,  
              struct sigaction *psOldAction);
```

- **iSig**: The type of signal to be affected
- **psAction**: Pointer to a structure containing instructions on how to handle signals of type **iSig**, including signal handler name and which signal types should be blocked
- **psOldAction**: (Irrelevant for our purposes)
- Installs an appropriate handler
- Automatically blocks signals of type **iSig**
- Returns 0 iff successful

Note: More powerful than C90 **signal()**

# Installing a Handler Example

Program testsigaction.c:

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>

static void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}
...
```

# Installing a Handler Example (cont.)

Program testsigaction.c (cont.):

```
...
int main(void) {
    int iRet;
    struct sigaction sAction;
    sAction.sa_flags = 0;
    sAction.sa_handler = myHandler;
    sigemptyset(&sAction.sa_mask);
    sigaction(SIGINT, &sAction, NULL);

    printf("Entering an infinite loop\n");
    for (;;)
        ;
    return 0;
}
```

# Alarm Example 2

Program testalarmtimeout.c:

```
#define _GNU_SOURCE
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <signal.h>
#include <unistd.h>

static void myHandler(int iSig)
{
    printf("\nSorry.  You took too long.\n");
    exit(EXIT_FAILURE);
}
```



# Interval Timers

**settimer()**

```
int settimer(int iWhich,  
             const struct itimerval *psValue,  
             struct itimerval *psOldValue);
```

- Sends SIGALRM
- Timing is specified by **psValue**
- **psOldValue** is irrelevant for our purposes
- Uses **virtual time**, alias **CPU time**
  - Time spent executing other processes does not count
  - Time spent waiting for user input does not count
- Returns 0 iff successful

# Interval Timer Example

Program testtimer.c:

```
##include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <signal.h>
#include <sys/time.h>

static void myHandler(int iSig) {
    printf("In myHandler with argument %d\n", iSig);
}
...
```

# Interval Timer Example (cont.)

Program testitimer.c (cont.):

```
...
int main(void)
{
    int iRet;
    void (*pfRet)(int);
    struct itimerval sTimer;

    pfRet = signal(SIGPROF, myHandler);
    assert(pfRet != SIG_ERR);
    ...
}
```

# Interval Timer Example (cont.)

Program testitimer.c (cont.):

```
    /* Send first signal in 1 second, 0 microseconds.
    */
    sTimer.it_value.tv_sec = 1;
    sTimer.it_value.tv_usec = 0;


    /* Send subsequent signals in 1 second,
       0 microseconds intervals. */
    sTimer.it_interval.tv_sec = 1;
    sTimer.it_interval.tv_usec = 0;

    iRet = setitimer(ITIMER_PROF, &sTimer, NULL);
    assert(iRet != -1);

    printf("Entering an infinite loop\n");
    for (;;)
        ;
    return 0;
}
```


# IPC using Pipes

- Communication channel between two processes. Like what happens when you type – `cat foo | grep bar`. The communication between the *stdout* of `foo` to the *stdin* of `grep`




# IPC using Pipes

- What happens when a process `fork()`s while still having pipes.



# IPC using Pipes

- What you really want is this:



# IPC Example: Unix pipe

- `int pipe(int fd[2]);`
  - Returns two file descriptors in `fd[0]` and `fd[1]`;
  - Writes to `fd[1]` will be read on `fd[0]`
  - When last copy of `fd[1]` closed, `fd[0]` will return EOF
  - Return 0 on success, -1 on error
- Operations on pipes:
  - read/write/close --- as with files
  - When `fd[1]` closed, `read(fd[0])` returns 0 bytes
  - When `fd[0]` closed, `write(fd[1])`:
    - Kills process with SIGPIPE, or if blocked
    - Fails with EPIPE



# IPC Example: Unix pipe (cont.)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(void)
{
    int pfd[2];
    pipe(pfd);
    if (!fork())
    {
        close(1); /* close normal stdout */
        dup(pfd[1]); /* make stdout same as pfd[1] */
        close(pfd[0]); /* we don't need this */
        execlp("ls", "ls", NULL); }
    else {
        close(0); /* close normal stdin */
        dup(pfd[0]); /* make stdin same as pfd[0] */
        close(pfd[1]); /* we don't need this */
        execlp("wc", "wc", "-l", NULL);
    }
    return 0; }
```

# Cooperating Processes

- ***Independent*** process cannot affect or be affected by the execution of another process
- ***Cooperating*** process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience

# Named Pipes – aka FIFOs

- System primitives used for implementing IPC when no two processes are related.
- Pipe()s work when the parents are directly related to one another – parent/child.
- What if the two processes are not related to one another.

# Named Pipes – aka FIFOs

- Creating a FIFO

```
mknod("myfifo", S_IFIFO | 0644 , 0);
```

``myfifo``: name of the FIFO

S\_IFIFO: Create a FIFO

S\_IFIFO|0644: add permissions : rw, r, r

# Named Pipes – aka FIFOs

```
#define FIFO_NAME "fifo_1"

int main(void)
{ char s[300];
  int num, fd;
  mknod(FIFO_NAME, S_IFIFO | 0666, 0);
  printf("waiting for readers...\n");
  fd = open(FIFO_NAME, O_WRONLY);
  printf("got a reader--type some stuff\n");
  while (gets(s), !feof(stdin))
  {
    if ((num = write(fd, s, strlen(s))) == -1)
      perror("write");
    else printf("speak: wrote %d bytes\n", num);
  }
  return 0;
}
```

```
#define FIFO_NAME "fifo_1"

int main(void)
{ char s[300];
  int num, fd;
  mknod(FIFO_NAME, S_IFIFO | 0666, 0);
  printf("waiting for writers...\n");
  fd = open(FIFO_NAME, O_WRONLY);
  printf("got a reader--type some stuff\n");
  do
  {
    if ((num = read(fd, s, 300)) == -1)
      perror("read");
    else
      { s[num] = '\0'; printf("read %d bytes: \"%s\"\n", num, s);
      }
  } while (num>0);
  return 0;
}
```

# Shared Memory – Memory section that can be shared between different processes

- Create a shared memory ID, much like a file descriptor.

```
int shmget(key_t key, size_t size, int shmflg);
```

```
key_t key;
```

```
int shmid;
```

```
key = ftok("/home/beej/somefile3", 'R');
```

```
shmid = shmget(key, 1024, 0644 | IPC_CREAT);
```

# Shared Memory – Attach to an address

- Attach to memory location or let the kernel decide one for you
- `void *shmat(int shmid, void *shmaddr, int shmflg);`

```
key_t key;  
int shmid;  
char *data;  
key = ftok("/home/beej/somefile3", 'R');  
shmid = shmget(key, 1024, 0644 | IPC_CREAT);  
data = shmat(shmid, (void *)0, 0);
```

You may read or write to data like you do with any memory location.

Once you are done with the shared memory you may need to detach it.

```
int shmdt(void *shmaddr);
```

But is not yet destroyed. To actually destroy it you need to actually delete it

```
shmctl(shmid, IPC_RMID, NULL);
```

# Shared Memory – with multiple processes

```
#define SHMSIZE 27
int main()
{
    int shmid;
    char *shm;
    if(fork() == 0)
    {
        shmid = shmget(2009, SHMSIZE, 0);
        shm = shmat(shmid, 0, 0);
        char *s = (char *) shm; *s = '\0';
        int i;
        for(i=0; i<5; i++)
        {
            int n;
            printf("Enter number<%i>: ", i);    scanf("%d", &n);
            sprintf(s, "%s%d", s, n);
        }
        strcat(s, "\n");
        printf ("Child wrote <%s>\n",shm);
        shmdt(shm);
    }
    else {
        shmid = shmget(2009, SHMSIZE, 0666 | IPC_CREAT);
        shm = shmat(shmid, 0, 0);
        wait(NULL);
        printf ("Parent reads <%s>\n",shm) ;
        shmdt(shm);
        shmctl(shmid, IPC_RMID, NULL);
    }
    return 0;
}
```