

# ***Operating Systems***

**CSE 231**

**Instructor: Sambuddho Chakravarty**

(Semester: Monsoon 2020)

Week 5: Oct 5 – Oct 8

# IPC Using Message Passing

- Very similar to FIFOs.
- Relies on message queues (again FIFO).

- Acquire message queue ID

```
int msgget(key_t key, int msgflg);
```

# IPC Using Message Passing

- Associate the message queue to an ``key'' much like other IPC mechanisms.

```
#include <sys/msg.h>
```

```
key = ftok("/home/beej/somefile", 'b');
```

```
msqid = msgget(key, 0666 | IPC_CREAT);
```

# IPC Using Message Passing

- Message buffer type (defined in sys/msg.h)

```
struct msgbuf  
{ long mtype;  
  char mtext[1]; };
```

Problem: mtext only single byte!

# IPC Using Message Passing

- You can use any structure as long as the `mtype` is type `long`.

```
struct pirate_msgbuf
{ long mtype; /* must be positive */
  struct pirate_info {
    char name[30];
    char ship_type;
    ..} info;
};
```

# IPC Using Message Passing

- Sending the message:

```
int msgsnd(int msqid, const void *msgp, size_t msgsz, int msgflg);
```

- Receiving the message:

```
int msgrcv(int msqid, void *msgp, size_t msgsz, long msgtyp, int msgflg);
```

# IPC Using Message Passing

- Msgtype used to determine which message to receive from the queue.

| <u><i>msgtyp</i></u> | <u>Effect on <i>msgrcv()</i></u>  |
|----------------------|---|
| Zero                 | Retrieve the next message on the queue, regardless of its <i>mtype</i> .  |
| Positive             | Get the next message with an <i>mtype equal</i> to the specified <i>msgtyp</i> .  |
| Negative             | Retrieve the first message on the queue whose <i>mtype</i> field is less than or equal to the absolute value of the <i>msgtyp</i> argument. |

# IPC Using Message Passing

- Delete the message queue:

```
int msgctl(int msqid, int cmd, struct msqid_ds *buf);
```

```
#include <sys/msg.h>
```

```
.
```

```
.
```

```
.
```

```
msgctl(msqid, IPC_RMID, NULL);
```



# IPC Using Unix Domain Sockets

- Domain sockets: Much like INET sockets used for communication between clients and servers.
- Two-way FIFOs.
- Socket structure

```
struct sockaddr_un
{
    unsigned short sun_family; /* AF_UNIX */
    char sun_path[108];
}
```

# IPC Using Unix Domain Sockets

- Steps to create a server socket:

socket() system call.

```
unsigned int s, s2;
```

```
struct sockaddr_un local, remote;
```

```
int len;
```

```
s = socket(AF_UNIX, SOCK_STREAM, 0);
```

# IPC Using Unix Domain Sockets

- Steps to create a server socket:

`bind()` system call – associates a “socket address” (a file path) to the said socket (much like what was done using shared memory)

```
local.sun_family = AF_UNIX; /* local is declared before socket() ^ */  
strcpy(local.sun_path, "/home/beej/mysocket");  
unlink(local.sun_path);  
len = strlen(local.sun_path) + sizeof(local.sun_family);  
bind(s, (struct sockaddr *)&local, len)
```

# IPC Using Unix Domain Sockets

- Steps to create a server socket:

`listen()` system call – Set the socket ``state'' to waiting.

```
listen(s, 5);
```

# IPC Using Unix Domain Sockets

- Steps to create a server socket:

`accept()` system call – Accept incoming connections from a client. Returns a connected socket descriptor (different from the previously created socket descriptor)

```
len = sizeof(struct sockaddr_un);
```

```
s2 = accept(s, &remote, &len); /* s2 is the newly connected  
socket descriptor which the server uses to communicate to  
the client*/
```

# IPC Using Unix Domain Sockets

- Steps to create a server socket:

send() or recv() to the connected socket descriptor.

```
while (len = recv(s2, &buf, 100, 0), len > 0)
    send(s2, &buf, len, 0);
```

```
/* loop back to accept() from here */
```

# IPC Using Unix Domain Sockets

- Steps to create a server socket:

When done with the communication, you `close()` or `shutdown()` the socket

`close(s);`

Or

`shutdown(s,how);`

`how == (SHUT_RD || SHUT_WR || SHUT_RDWR)`

# IPC Using Unix Domain Sockets

- Client part of the communication
- Initiate a client socket()

```
int s;
```

```
struct sockaddr_un remote;
```

```
if ((s = socket(AF_UNIX, SOCK_STREAM, 0)) == -1)  
{  
    perror("socket"); exit(1);  
}
```



# IPC Using Unix Domain Sockets

- Connect() to the server

```
#define SOCK_PATH "/path/to/mysocket"
```

```
struct sockaddr_un remote;
```

```
remote.sun_family = AF_UNIX;
```

```
strcpy(remote.sun_path, SOCK_PATH);
```

```
len = strlen(remote.sun_path) + sizeof(remote.sun_family);
```

```
if (connect(s, (struct sockaddr *)&remote, len) == -1)
```

```
{ perror("connect"); exit(1);
```

```
}
```

# IPC Using Unix Domain Sockets

- `send()` / `recv()` to/from the socket just like you do for the server socket

# Socketpair – Full-duplex pipe (socket + pipe)

```
int main(void)
{ int sv[2]; /* the pair of socket descriptors */
  char buf; /* for data exchange between processes */
  if (socketpair(AF_UNIX, SOCK_STREAM, 0, sv) == -1)
  { perror("socketpair"); exit(1); }
  if (!fork())
  { read(sv[1], &buf, 1);
    printf("child: read '%c'\n", buf);
    buf = toupper(buf); /* make it uppercase */
    write(sv[1], &buf, 1);
    printf("child: sent '%c'\n", buf); }
  else { /* parent */ write(sv[0], "b", 1);
    printf("parent: sent 'b'\n");
    read(sv[0], &buf, 1); p
    printf("parent: read '%c'\n", buf);
    wait(NULL); /* wait for child to die */ }
  return 0;}
```